

[android](#) / [platform](#) / [external](#) / [qemu](#) / [refs/heads/main](#) / [.](#) / [docs](#) / **GOLDFISH-VIRTUAL-HARDWARE.TXT**

blob: 3e410888ef9e558b10ad7528bac221b9c029e290 [[file](#)] [[log](#)] [[blame](#)]

```
1 Introduction
2 =====
3
4 This file documents the 'goldfish' virtual hardware platform used to run some
5 emulated Android systems under QEMU. It serves as a reference for implementers
6 of virtual devices in QEMU, as well as Linux kernel developers who need to
7 maintain the corresponding drivers.
8
9 The following abbreviations will be used here:
10
11     $QEMU -> path to the Android AOSP directory, i.e. a git clone of
12             https://android.googlesource.com/platform/external/qemu.git
13
14     $KERNEL -> path to the Android goldfish kernel source tree, i.e. a git clone of
15                https://android.googlesource.com/kernel/goldfish.git
16
17             More specifically, to the android-goldfish-2.6.29 branch for now.
18
19 'goldfish' is the name of a family of similar virtual hardware platforms, that
20 mostly differ in the virtual CPU they support. 'goldfish' started as an
21 ARM-specific platform, but has now been ported to x86 and MIPS virtual CPUs.
22
23 Inside of QEMU, goldfish-specific virtual device implementation sources files
24 are in $QEMU/hw/android/goldfish/*.c sources
25
26 Inside the Linux kernel tree, they are under $KERNEL/arch/$ARCH/mach-goldfish,
27 or $KERNEL/arch/$ARCH/goldfish/, as well as a few arch-independent drivers in
28 different locations (detailed below).
29
30 Goldfish devices appear to the Linux kernel as 'platform devices'. Read [1] and
31 [2] for an introduction and reference documentation for these.
32
33 Each device is identified by a name, and an optional unique id (an integer used
34 to distinguish between several identical instances of similar devices, like
35 serial ports, of block devices). When only one instance of a given device can
36 be used, an ID of -1 is used instead.
37
38 It also communicates with the kernel through:
39
40     - One or more 32-bit of I/O registers, mapped to physical addresses at
41       specific locations which depend on the architecture.
42
43     - Zero or more interrupt requests, used to signal to the kernel that an
44       important event occurred.
45
46     Note that IRQ lines are numbered from 0 to 31, and are relative to the
```

```
47     goldfish interrupt controller, documented below.
48
49
50 [1] http://lwn.net/Articles/448499/
51 [2] https://www.kernel.org/doc/Documentation/driver-model/platform.txt
52
53
54 I. Goldfish platform bus:
55 =====
56
57 The 'platform bus', in Linux kernel speak, is a special device that is capable
58 of enumerating other platform devices found on the system to the kernel. This
59 flexibility allows to customize which virtual devices are available when running
60 a given emulated system configuration.
61
62 Relevant files:
63     $QEMU/hw/android/goldfish/device.c
64     $KERNEL/arch/arm/mach-goldfish/pdev_bus.c
65     $KERNEL/arch/x86/mach-goldfish/pdev_bus.c
66     $KERNEL/arch/mips/goldfish/pdev_bus.c
67
68 Device properties:
69     Name: goldfish_device_bus
70     Id:   -1
71     IrqCount: 1
72
73 32-bit I/O registers (offset, name, abstract)
74
75     0x00 BUS_OP      R: Iterate to next device in enumeration.
76                     W: Start device enumeration.
77
78     0x04 GET_NAME    W: Copy device name to kernel memory.
79     0x08 NAME_LEN    R: Read length of current device's name.
80     0x0c ID          R: Read id of current device.
81     0x10 IO_BASE     R: Read I/O base address of current device.
82     0x14 IO_SIZE     R: Read I/O base size of current device.
83     0x18 IRQ_BASE    R: Read base IRQ of current device.
84     0x1c IRQ_COUNT   R: Read IRQ count of current device.
85
86     # For 64-bit guest architectures only:
87     0x20 NAME_ADDR_HIGH W: Write high 32-bit of kernel address of name
88                         buffer used by GET_NAME. Must be written to
89                         before the GET_NAME write.
90
91 The kernel iterates over the list of current devices with something like:
92
93     IO_WRITE(BUS_OP, 0);    // Start iteration, any value other than 0 is invalid.
94     for (;;) {
95         int ret = IO_READ(BUS_OP);
96         if (ret == 0 /* OP_DONE */) {
97             // no more devices.
98             break;
99         }
100     else if (ret == 8 /* OP_ADD_DEV */) {
101         // Read device properties.
```

```

102     Device dev;
103     dev.name_len  = IO_READ(NAME_LEN);
104     dev.id        = IO_READ(ID);
105     dev.io_base   = IO_READ(IO_BASE);
106     dev.io_size   = IO_READ(IO_SIZE);
107     dev.irq_base  = IO_READ(IRQ_BASE);
108     dev.irq_count = IO_READ(IRQ_COUNT);
109
110     dev.name = kalloc(dev.name_len + 1); // allocate room for device name.
111     #if 64BIT_GUEST_CPU
112         IO_WRITE(NAME_ADDR_HIGH, (uint32_t)(dev.name >> 32));
113     #endif
114     IO_WRITE(GET_NAME, (uint32_t)dev.name); // copy to kernel memory.
115     dev.name[dev.name_len] = 0;
116
117     .. add device to kernel's list.
118 }
119 else {
120     // Not returned by current goldfish implementation.
121 }
122 }
123
124 The device also uses a single IRQ, which it will raise to indicate to the kernel
125 that new devices are available, or that some of them have been removed. The
126 kernel will then start a new enumeration. The IRQ is lowered by the device only
127 when a IO_READ(BUS_OP) returns 0 (OP_DONE).
128
129 NOTE: The kernel hard-codes a platform_device definition with the name
130 "goldfish_pdev_bus" for the platform bus (e.g. see
131 $KERNEL/arch/arm/mach-goldfish/board-goldfish.c), however, the bus itself
132 will appear during enumeration as a device named "goldfish_device_bus"
133
134 The kernel driver for the platform bus only matches the "goldfish_pdev_bus"
135 name, and will ignore any device named "goldfish_device_bus".
136
137
138 II. Goldfish interrupt controller:
139 =====
140
141 IMPORTANT: The controller IS NOT USED IN EMULATED X86 SYSTEMS.
142     TODO(digit): Indicate which virtual PIC is used on x86 systems.
143
144 Relevant files:
145     $QEMU/hw/android/goldfish/interrupt.c
146     $KERNEL/arch/arm/mach-goldfish/board-goldfish.c
147     $KERNEL/arch/mips/goldfish/goldfish-interrupt.c
148
149 Device properties:
150     Name: goldfish_interrupt_controller
151     Id: -1
152     IrqCount: 0 (uses parent CPU IRQ instead).
153
154     32-bit I/O registers (offset, name, abstract):
155     0x00 STATUS      R: Read the number of pending interrupts (0 to 32).
156     0x04 NUMBER      R: Read the lowest pending interrupt index, or 0 if none.

```

```

157     0x08 DISABLE_ALL    W: Clear all pending interrupts (does not disable them!)
158     0x0c DISABLE       W: Disable a given interrupt, value must be in [0..31].
159     0x10 ENABLE        W: Enable a given interrupt, value must be in [0..31].
160
161 Goldfish provides its own interrupt controller that can manage up to 32 distinct
162 maskable interrupt request lines. The controller itself is cascaded from a
163 parent CPU IRQ.
164
165 What this means in practice:
166
167     - Each IRQ has a 'level' that is either 'high' (1) or 'low' (0).
168
169     - Each IRQ also has a binary 'enable' flag.
170
171     - Whenever (level == 1 && enabled == 1) is reached due to a state change, the
172       controller raises its parent IRQ. This typically interrupts the CPU and
173       forces the kernel to deal with the interrupt request.
174
175     - Raised/Enabled interrupts that have not been serviced yet are called
176       "pending". Raised/Disabled interrupts are called "masked" and are
177       essentially silent until enabled.
178
179 When the interrupt controller triggers the parent IRQ, the kernel should do
180 the following:
181
182     num_pending = IO_READ(STATUS); // Read number of pending interrupts.
183     for (int n = 0; n < num_pending; ++n) {
184         int irq_index = IO_READ(NUMBER); // Read n-th interrupt index.
185         .. service interrupt request with the proper driver.
186     }
187
188 IO_WRITE(DISABLE, <num>) or IO_WRITE(ENABLE, <num>) can change the 'enable' flag
189 of a given IRQ. <num> must be a value in the [0..31] range. Note that enabling
190 an IRQ which has already been raised will make it active, i.e. it will raise
191 the parent IRQ.
192
193 IO_WRITE(DISABLE_ALL, 0) can be used to lower all interrupt levels at once (even
194 disabled one). Note that this constant is probably mis-named since it does not
195 change the 'enable' flag of any IRQ.
196
197 Note that this is the only way for the kernel to lower an IRQ level through
198 this device. Generally speaking, Goldfish devices are responsible for lowering
199 their own IRQ, which is performed either when a specific condition is met, or
200 when the kernel reads from or writes to a device-specific I/O register.
201
202
203 III. Godlfish timer:
204 =====
205
206 NOTE: This is not used on x86 emulated platforms.
207
208 Relevant files:
209     $QEMU/hw/android/goldfish/timer.c
210     $KERNEL/arch/arm/mach-goldfish/timer.c
211     $KERNEL/arch/mips/goldfish/goldfish-time.c

```

```

212
213 Device properties:
214     Name: goldfish_timer
215     Id: -1
216     IrqCount: 1
217
218     32-bit I/O registers (offset, name, abstract)
219         0x00  TIME_LOW           R: Get current time, then return low-order 32-bits.
220         0x04  TIME_HIGH          R: Return high 32-bits from previous TIME_LOW read.
221         0x08  ALARM_LOW          W: Set low 32-bit value of alarm, then arm it.
222         0x0c  ALARM_HIGH         W: Set high 32-bit value of alarm.
223         0x10  CLEAR_INTERRUPT    W: Lower device's irq level.
224         0x14  CLEAR_ALARM
225
226 This device is used to return the current host time to the kernel, as a
227 high-precision signed 64-bit nanoseconds value, starting from a liberal point
228 in time. This value should correspond to the QEMU "vm_clock", i.e. it should
229 not be updated when the emulated system does _not_ run, and hence cannot be
230 based directly on a host clock.
231
232 To read the value, the kernel must perform an IO_READ(TIME_LOW), which returns
233 an unsigned 32-bit value, before an IO_READ(TIME_HIGH), which returns a signed
234 32-bit value, corresponding to the higher half of the full value.
235
236 The device can also be used to program an alarm, with something like:
237
238     IO_WRITE(ALARM_HIGH, <high-value>) // Must happen first.
239     IO_WRITE(ALARM_LOW, <low-value>)   // Must happen second.
240
241 When the corresponding value is reached, the device will raise its IRQ. Note
242 that the IRQ is raised as soon as the second IO_WRITE() if the alarm value is
243 already older than the current time.
244
245 IO_WRITE(CLEAR_INTERRUPT, <any>) can be used to lower the IRQ level once the
246 alarm has been handled by the kernel.
247
248 IO_WRITE(CLEAR_ALARM, <any>) can be used to disarm an existing alarm, if any.
249
250 Note: At the moment, the alarm is only used on ARM-based system. MIPS based
251 systems only use TIME_LOW / TIME_HIGH on this device.
252
253
254 III. Goldfish real-time clock (RTC):
255 =====
256
257 Relevant files:
258     $QEMU/hw/android/goldfish/timer.c
259     $KERNEL/drivers/rtc/rtc-goldfish.c
260
261 Device properties:
262     Name: goldfish_rtc
263     Id: -1
264     IrqCount: 1
265     I/O Registers:
266         0x00  TIME_LOW           R: Get current time, then return low-order 32-bits.

```

```

267     0x04  TIME_HIGH      R: Return high 32-bits, from previous TIME_LOW read.
268     0x08  ALARM_LOW     W: Set low 32-bit value or alarm, then arm it.
269     0x0c  ALARM_HIGH    W: Set high 32-bit value of alarm.
270     0x10  CLEAR_INTERRUPT W: Lower device's irq level.

```

271

272 This device is `_very_` similar to the Goldfish timer one, with the following
 273 important differences:

274

- 275 - Values reported are still 64-bit nanoseconds, but they have a granularity
 276 of 1 second, and represent host-specific values (really `'time() * 1e9'`)
- 277
- 278 - The alarm is non-functioning, i.e. writing to `ALARM_LOW` / `ALARM_HIGH` will
 279 work, but will never arm any alarm.

280

281 To support old Goldfish kernels, make sure to support writing to
 282 `ALARM_LOW` / `ALARM_HIGH` / `CLEAR_INTERRUPT`, even if the device never raises its
 283 IRQ.

284

285

286 IV. Goldfish serial port (tty):

287 =====

288

289 Relevant files:

```

290     $QEMU/hw/android/goldfish/tty.c
291     $KERNEL/drivers/char/goldfish_tty.c
292     $KERNEL/arch/arm/mach-goldfish/include/debug-macro.S

```

293

294 Device properties:

295 Name: goldfish_tty

296 Id: 0 to N

297 IrqCount:

298 I/O Registers:

```

299     0x00  PUT_CHAR      W: Write a single 8-bit value to the serial port.
300     0x04  BYTES_READY   R: Read the number of available buffered input bytes.
301     0x08  CMD           W: Send command (see below).
302     0x10  DATA_PTR     W: Write kernel buffer address.
303     0x14  DATA_LEN     W: Write kernel buffer size.

```

304

305 # For 64-bit guest CPUs only:

```

306     0x18  DATA_PTR_HIGH W: Write high 32 bits of kernel buffer address.

```

307

308 This is the first case of a multi-instance goldfish device in this document.
 309 Each instance implements a virtual serial port that contains a small internal
 310 buffer where incoming data is stored until the kernel fetches it.

311

312 The CMD I/O register is used to send various commands to the device, identified
 313 by the following values:

314

```

315     0x00  CMD_INT_DISABLE  Disable device.
316     0x01  CMD_INT_ENABLE   Enable device.
317     0x02  CMD_WRITE_BUFFER Write buffer from kernel to device.
318     0x03  CMD_READ_BUFFER  Read buffer from device to kernel.

```

319

320 Each device instance uses one IRQ that is raised to indicate that there is
 321 incoming/buffered data to read. To read such data, the kernel should do the

following:

```

len = IO_READ(PUT_CHAR);    // Read length of incoming data.
if (len == 0) return;       // Nothing to do.

available = get_buffer(len, &buffer); // Get address of buffer and its size.
#ifdef 64BIT_GUEST_CPU
IO_WRITE(DATA_PTR_HIGH, buffer >> 32);
#endif
IO_WRITE(DATA_PTR, buffer);    // Write buffer address to device.
IO_WRITE(DATA_LEN, available); // Write buffer length to device.
IO_WRITE(CMD, CMD_READ_BUFFER); // Read the data into kernel buffer.

```

The device will automatically lower its IRQ when there is no more input data in its buffer. However, the kernel can also temporarily disable device interrupts with `CMD_INT_DISABLE / CMD_INT_ENABLE`.

Note that disabling interrupts does not flush the buffer, nor prevent it from buffering further data from external inputs.

To write to the serial port, the device can either send a single byte at a time with:

```
IO_WRITE(PUT_CHAR, <value>)    // Send the lower 8 bits of <value>.
```

Or use the more efficient sequence:

```

#ifdef 64BIT_GUEST_CPU
IO_WRITE(DATA_PTR_HIGH, buffer >> 32)
#endif
IO_WRITE(DATA_PTR, buffer)
IO_WRITE(DATA_LEN, buffer_len)
IO_WRITE(CMD, CMD_WRITE_BUFFER)

```

The former is less efficient but simpler, and is typically used by the kernel to send debug messages only.

Note that the Android emulator always reserves the first two virtual serial ports:

- The first one is used to receive kernel messages, this is done by adding the 'console=ttyS0' parameter to the kernel command line in `$QEMU/vl-android.c`
- The second one is used to setup the legacy "qemud" channel, used on older Android platform revisions. This is done by adding 'android.qemud=ttyS1' on the kernel command line in `$QEMU/vl-android.c`

Read `docs/ANDROID-QEMUD.TXT` for more details about the data that passes through this serial port. In a nutshell, this is required to emulate older Android releases (e.g. cupcake). It provides a direct communication channel between the guest system and the emulator.

More recent Android platforms do not use QEMUD anymore, but instead rely on the much faster "QEMU pipe" device, described later in this document as

```
377     well as in docs/ANDROID-QEMU-PIPE.TXT.
378
379
380 V. Goldfish framebuffer:
381 =====
382
383 Relevant files:
384     $QEMU/hw/android/goldfish/fb.c
385     $KERNEL/drivers/video/goldfish_fb.c
386
387 Device properties:
388     Name: goldfish_fb
389     Id: 0 to N (only one used in practice).
390     IrqCount: 0
391     I/O Registers:
392         0x00  GET_WIDTH      R: Read framebuffer width in pixels.
393         0x04  GET_HEIGHT    R: Read framebuffer height in pixels.
394         0x08  INT_STATUS
395         0x0c  INT_ENABLE
396         0x10  SET_BASE
397         0x14  SET_ROTATION
398         0x18  SET_BLANK      W: Set 'blank' flag.
399         0x1c  GET_PHYS_WIDTH R: Read framebuffer width in millimeters.
400         0x20  GET_PHYS_HEIGHT R: Read framebuffer height in millimeters.
401         0x24  GET_FORMAT    R: Read framebuffer pixel format.
402
403 The framebuffer device is a bit peculiar, because it uses, in addition to the
404 typical I/O registers and IRQs, a large area of physical memory, allocated by
405 the kernel, but visible to the emulator, to store a large pixel buffer.
406
407 The emulator is responsible for displaying the framebuffer content in its UI
408 window, which can be rotated, as instructed by the kernel.
409
410 IMPORTANT NOTE: When GPU emulation is enabled, the framebuffer will typically
411 only be used during boot. Note that GPU emulation doesn't rely on a specific
412 virtual GPU device, however, it uses the "QEMU Pipe" device described below.
413 For more information, please read:
414
415     external/qemu/distrib/android-emugl/DESIGN
416
417 On boot, the kernel will read various properties of the framebuffer:
418
419     IO_READ(GET_WIDTH) and IO_READ(GET_HEIGHT) return the width and height of
420     the framebuffer in pixels. Note that a 'row' corresponds to consecutive bytes
421     in memory, but doesn't necessarily to an horizontal line on the final display,
422     due to possible rotation (see SET_ROTATION below).
423
424     IO_READ(GET_PHYS_WIDTH) and IO_READ(GET_PHYS_HEIGHT) return the emulated
425     physical width and height in millimeters, this is later used by the kernel
426     and the platform to determine the device's emulated density.
427
428     IO_READ(GET_FORMAT) returns a value matching the format of pixels in the
429     framebuffer. Note that these values are specified by the Android hardware
430     abstraction layer (HAL) and cannot change:
431
```



```
432     0x01  HAL_PIXEL_FORMAT_BGRA_8888
433     0x02  HAL_PIXEL_FORMAT_RGBX_8888
434     0x03  HAL_PIXEL_FORMAT_RGB_888
435     0x04  HAL_PIXEL_FORMAT_RGB_565
436     0x05  HAL_PIXEL_FORMAT_BGRA_8888
437     0x06  HAL_PIXEL_FORMAT_RGBA_5551
438     0x08  HAL_PIXEL_FORMAT_RGBA_4444
439
440     HOWEVER, the kernel driver only expects a value of HAL_PIXEL_FORMAT_RGB_565
441     at the moment. Until this is fixed, the virtual device should always return
442     the value 0x04 here. Rows are not padded, so the size in bytes of a single
443     framebuffer will always be exactly 'width * height * 2'.
444
445     Note that GPU emulation doesn't have this limitation and can use and display
446     32-bit surfaces properly, because it doesn't use the framebuffer.
447
448     The device has a 'blank' flag. When set to 1, the UI should only display an
449     empty/blank framebuffer, ignoring the content of the framebuffer memory.
450     It is set with IO_WRITE(SET_BLANK, <value>), where value can be 1 or 0. This is
451     used when simulating suspend/resume.
452
453     IMPORTANT: The framebuffer memory is allocated by the kernel, which will send
454     its physical address to the device by using IO_WRITE(SET_BASE, <address>).
455
456     The kernel really allocates a memory buffer large enough to hold *two*
457     framebuffers, in order to implement panning / double-buffering. This also means
458     that calls to IO_WRITE(SET_BASE, <address>) will be frequent.
459
460     The allocation happens with dma_alloc_writecombine() on ARM, which can only
461     allocate a maximum of 4 MB, this limits the size of each framebuffer to 2 MB,
462     which may not be enough to emulate high-density devices :-(
463
464     For other architectures, dma_alloc_coherent() is used instead, and has the same
465     upper limit / limitation.
466
467     TODO(digit): Explain how it's possible to raise this limit by modifyinf
468                  CONSISTENT_DMA_SIZE and/or MAX_ORDER in the kernel configuration.
469
470     The device uses a single IRQ to notify the kernel of several events. When it
471     is raised, the kernel IRQ handler must IO_READ(INT_STATUS), which will return
472     a value containing the following bit flags:
473
474         bit 0: Set to 1 to indicate a VSYNC event.
475
476         bit 1: Set to 1 to indicate that the content of a previous SET_BASE has
477                been properly displayed.
478
479     Note that reading this register also lowers the device's IRQ level.
480
481     The second flag is essentially a way to notify the kernel that an
482     IO_WRITE(SET_BASE, <address>) operation has been succesfully processed by
483     the emulator, i.e. that the new content has been displayed to the user.
484
485     The kernel can control which flags should raise an IRQ by using
486     IO_WRITE(INT_ENABLE, <flags>), where <flags> has the same format as the
```

```

487 result of IO_READ(INT_STATUS). If the corresponding bit is 0, the an IRQ
488 for the corresponding event will never be generated,
489
490
491 VI. Goldfish audio device:
492 =====
493
494 Relevant files:
495     $QEMU/hw/android/goldfish/audio.c
496     $KERNEL/drivers/misc/goldfish_audio.c
497
498 Device properties:
499     Name: goldfish_audio
500     Id: -1
501     IrqCount: 1
502     I/O Registers:
503         0x00  INT_STATUS
504         0x04  INT_ENABLE
505         0x08  SET_WRITE_BUFFER_1      W: Set address of first kernel output buffer.
506         0x0c  SET_WRITE_BUFFER_2      W: Set address of second kernel output buffer.
507         0x10  WRITE_BUFFER_1          W: Send first kernel buffer samples to output.
508         0x14  WRITE_BUFFER_2          W: Send second kernel buffer samples to output.
509         0x18  READ_SUPPORTED          R: Reads 1 if input is supported, 0 otherwise.
510         0x1c  SET_READ_BUFFER
511         0x20  START_READ
512         0x24  READ_BUFFER_AVAILABLE
513
514     # For 64-bit guest CPUs
515         0x28  SET_WRITE_BUFFER_1_HIGH W: Set high 32 bits of 1st kernel output buffer address.
516         0x30  SET_WRITE_BUFFER_2_HIGH W: Set high 32 bits of 2nd kernel output buffer address.
517         0x34  SET_READ_BUFFER_HIGH    W: Set high 32 bits of kernel input buffer address.
518
519 This device implements a virtual sound card with the following properties:
520
521     - Stereo output at fixed 44.1 kHz frequency, using signed 16-bit samples.
522       Mandatory.
523
524     - Mono input at fixed 8 kHz frequency, using signed 16-bit samples.
525       Optional.
526
527 For output, the kernel driver allocates two internal buffers to hold output
528 samples, and passes their physical address to the emulator as follows:
529
530     #if 64BIT_GUEST_CPU
531     IO_WRITE(SET_WRITE_BUFFER_1_HIGH, (uint32_t)(buffer1 >> 32));
532     IO_WRITE(SET_WRITE_BUFFER_2_HIGH, (uint32_t)(buffer2 >> 32));
533     #endif
534     IO_WRITE(WRITE_BUFFER_1, (uint32_t)buffer1);
535     IO_WRITE(WRITE_BUFFER_2, (uint32_t)buffer2);
536
537 After this, samples will be sent from the driver to the virtual device by
538 using one of IO_WRITE(WRITE_BUFFER_1, <length1>) or
539 IO_WRITE(WRITE_BUFFER_2, <length2>), depending on which sample buffer to use.
540 NOTE: Each length is in bytes.
541

```

542 Note however that the driver should wait, before doing this, until the device
543 gives permission by raising its IRQ and setting the appropriate 'status' flags.
544
545 The virtual device has an internal 'int_status' field made of 3 bit flags:
546
547 bit0: 1 iff the device is ready to receive data from the first buffer.
548 bit1: 1 iff the device is ready to receive data from the second buffer.
549 bit2: 1 iff the device has input samples for the kernel to read.
550
551 Note that an IO_READ(INT_STATUS) also automatically lowers the IRQ level,
552 except if the read value is 0 (which should not happen, since it should not
553 raise the IRQ).
554
555 The corresponding interrupts can be masked by using IO_WRITE(INT_ENABLE, <mask>),
556 where <mask> has the same format as 'int_status'. A 1 bit in the mask enables the
557 IRQ raise when the corresponding status bit is also set to 1.
558
559 For input, the driver should first IO_READ(READ_SUPPORTED), which will return 1
560 if the virtual device supports input, or 0 otherwise. If it does support it,
561 the driver must allocate an internal buffer and send its physical address with
562 IO_WRITE(SET_READ_BUFFER, <read-buffer>) (with a previous write to
563 SET_READ_BUFFER_HIGH on 64-bit guest CPUs), then perform
564 IO_WRITE(START_READ, <read-buffer-length>) to start recording and
565 specify the kernel's buffer length.
566
567 Later, the device will raise its IRQ and set bit2 of 'int_status' to indicate
568 there are incoming samples to the driver. In its interrupt handler, the latter
569 should IO_READ(READ_BUFFER_AVAILABLE), which triggers the transfer (from the
570 device to the kernel), as well as return the size in bytes of the samples.
571
572
573 VII. Goldfish battery:
574 =====
575
576 Relevant files:
577 \$QEMU/hw/android/goldfish/battery.c
578 \$QEMU/hw/power_supply.h
579 \$KERNEL/drivers/power/goldfish_battery.c
580
581 Device properties:
582 Name: goldfish_battery
583 Id: -1
584 IrqCount: 1
585 I/O Registers:
586 0x00 INT_STATUS R: Read battery and A/C status change bits.
587 0x04 INT_ENABLE W: Enable or disable IRQ on status change.
588 0x08 AC_ONLINE R: Read 0 if AC power disconnected, 1 otherwise.
589 0x0c STATUS R: Read battery status (charging/full/... see below).
590 0x10 HEALTH R: Read battery health (good/overheat/... see below).
591 0x14 PRESENT R: Read 1 if battery is present, 0 otherwise.
592 0x18 CAPACITY R: Read battery charge percentage in [0..100] range.
593
594 A simple device used to report the state of the virtual device's battery, and
595 whether the device is powered through a USB or A/C adapter.
596

```

597 The device uses a single IRQ to notify the kernel that the battery or A/C status
598 changed. When this happens, the kernel should perform an IO_READ(INT_STATUS)
599 which returns a 2-bit value containing flags:
600
601     bit 0: Set to 1 to indicate a change in battery status.
602     bit 1: Set to 1 to indicate a change in A/C status.
603
604 Note that reading this register also lowers the IRQ level.
605
606 The A/C status can be read with IO_READ(AC_ONLINE), which returns 1 if the
607 device is powered, or 0 otherwise.
608
609 The battery status is spread over multiple I/O registers:
610
611     IO_READ(PRESENT) returns 1 if the battery is present in the virtual device,
612     or 0 otherwise.
613
614     IO_READ(CAPACITY) returns the battery's charge percentage, as an integer
615     between 0 and 100, inclusive. NOTE: This register is probably misnamed since
616     it does not represent the battery's capacity, but it's current charge level.
617
618     IO_READ(STATUS) returns one of the following values:
619
620         0x00 UNKNOWN      Battery state is unknown.
621         0x01 CHARGING     Battery is charging.
622         0x02 DISCHARGING  Battery is discharging.
623         0x03 NOT_CHARGING Battery is not charging (e.g. full or dead).
624
625     IO_READ(HEALTH) returns one of the following values:
626
627         0x00 UNKNOWN      Battery health unknown.
628         0x01 GOOD         Battery is in good condition.
629         0x02 OVERHEATING  Battery is over-heating.
630         0x03 DEAD         Battery is dead.
631         0x04 OVERVOLTAGE  Battery generates too much voltage.
632         0x05 UNSPEC_FAILURE Battery has unspecified failure.
633
634 The kernel can use IO_WRITE(INT_ENABLE, <flags>) to select which condition
635 changes should trigger an IRQ. <flags> is a 2-bit value using the same format
636 as INT_STATUS.
637
638
639 VIII. Goldfish events device (user input):
640 =====
641
642 Relevant files:
643     $QEMU/hw/android/goldfish/events_device.c
644     $KERNEL/drivers/input/keyboard/goldfish_events.c
645
646 Device properties:
647     Name: goldfish_events
648     Id: -1
649     IrqCount: 1
650     I/O Registers:
651         0x00 READ      R: Read next event type, code or value.

```

```
652      0x00 SET_PAGE    W: Set page index.
653      0x04 LEN         R: Read length of page data.
654      0x08 DATA       R: Read page data.
655      ....            R: Read additional page data (see below).
```

656

657 This device is responsible for sending several kinds of user input events to
658 the kernel, i.e. emulated device buttons, hardware keyboard, touch screen,
659 trackball and lid events.

660

661 NOTE: Android supports other input devices like mice or game controllers
662 through USB or Bluetooth, these are not supported by this virtual
663 Goldfish device.

664

665 NOTE: The 'lid event' is useful for devices with a clamshell or foldable
666 keyboard design, and is used to report when it is opened or closed.

667

668 As per Linux conventions, each 'emulated event' is sent to the kernel as a
669 series of (<type>,<code>,<value>) triplets or 32-bit values. For more
670 information, see:

671

672 <https://www.kernel.org/doc/Documentation/input/input.txt>

673

674 As well as the <linux/input.h> kernel header.

675

676 Note that in the context of goldfish:

677

678 - Button and keyboard events are reported with:

679 (EV_KEY, <code>, <press>)

680

681 Where <code> is a 9-bit keycode, as defined by <linux/input.h>, and
682 <press> is 1 for key/button presses, and 0 for releases.

683

684 - For touchscreen events, a single-touch event is reported with:

685 (EV_ABS, ABS_X, <x-position>) +

686 (EV_ABS, ABS_Y, <y-position>) +

687 (EV_ABS, ABS_Z, 0) +

688 (EV_KEY, BTN_TOUCH, <button-state>) +

689 (EV_SYN, 0, 0)

690

691 where <x-position> and <y-position> are the horizontal and vertical position
692 of the touch event, respectfully, and <button-state> is either 1 or 0 and
693 indicates the start/end of the touch gesture, respectively.

694

695 - For multi-touch events, things are much more complicated. In a nutshell,
696 these events are reported through (EV_ABS, ABS_MT_XXXXX, YYY) triplets,
697 as documented at:

698

699 <https://www.kernel.org/doc/Documentation/input/multi-touch-protocol.txt>

700

701 TODO(digit): There may be bugs in either the virtual device or driver code
702 when it comes to multi-touch. Iron out the situation and better
703 explain what's required to support all Android platforms.

704

705 - For trackball events:

706 (EV_REL, REL_X, <x-delta>) +

```

707     (EV_REL, REL_Y, <y-delta>) +
708     (EV_SYN, 0, 0)
709
710     Where <x-delta> and <y-delta> are the signed relative trackball displacement
711     in the horizontal and vertical directions, respectively.
712
713     - For lid events:
714         (EV_SW, 0, 1) + (EV_SYN, 0, 0)    // When lid is closed.
715         (EV_SW, 0, 0) + (EV_SYN, 0, 0)    // When lid is opened.
716
717     When the kernel driver starts, it will probe the device to know what kind
718     of events are supported by the emulated configuration. There are several
719     categories of queries:
720
721     - Asking for the current physical keyboard 'charmap' name, used by the system
722       to translate keycodes in actual characters. In practice, this will nearly
723       always be 'goldfish' for emulated systems, but this out of spec for this
724       document.
725
726     - Asking which event codes are supported for a given event type
727       (e.g. all the possible KEY_XXX values generated for EV_KEY typed triplets).
728
729     - Asking for various minimum or maximum values for each supported EV_ABS
730       event code. For example the min/max values of (EV_ABS, ABS_X, ...) triplets,
731       to know the bounds of the input touch panel.
732
733     The kernel driver first select which kind of query it wants by using
734     IO_WRITE(SET_PAGE, <page>), where <page> is one of the following values:
735
736     PAGE_NAME      0x0000    Keyboard charmap name.
737     PAGE_EVBITS    0x10000    Event code supported sets.
738     PAGE_ABSDATA   0x20003    (really 0x20000 + EV_ABS) EV_ABS min/max values.
739
740     Once a 'page' has been selected, it is possible to read from it with
741     IO_READ(LEN) and IO_READ(DATA). In practice:
742
743     - To read the name of the keyboard charmap, the kernel will do:
744
745         IO_WRITE(SET_PAGE, PAGE_NAME); # Ind
746
747         charmap_name_len = IO_READ(LEN);
748         charmap_name = kalloc(charmap_name_len + 1);
749         for (int n = 0; n < charmap_name_len; ++n)
750             charmap_name[n] = (char) IO_READ(DATA);
751         charmap_name[n] = 0;
752
753     - To read which codes a given event type (here EV_KEY) supports:
754
755         IO_WRITE(SET_PAGE, PAGE_EVBITS + EV_KEY); // Or EV_REL, EV_ABS, etc...
756
757         bitmask_len = IO_READ(LEN);
758         for (int offset = 0; offset < bitmask_len; ++offset) {
759             uint8_t mask = (uint8_t) IO_READ(DATA);
760             for (int bit = 0; bit < 8; ++bit) {
761                 int code = (offset * 8) + bit;

```

```
762         if ((mask & (1 << bit)) != 0) {
763             ... record that keycode |code| is supported.
764         }
765     }
766 }
```

- To read the range values of absolute event values:

```
769
770     IO_WRITE(SET_PAGE, PAGE_ABSDATA);
771     max_entries = IO_READ(LEN);
772     for (int n = 0; n < max_entries; n += 4) {
773         int32_t min = IO_READ(DATA + n);
774         int32_t max = IO_READ(DATA + n + 4);
775         int32_t fuzz = IO_READ(DATA + n + 8);
776         int32_t flat = IO_READ(DATA + n + 12);
777         int event_code = n/4;
778
779         // Record (min, max, fuzz, flat) values for EV_ABS 'event_code'.
780     }
781
```

Note that the 'fuzz' and 'flat' values reported by Goldfish are always 0, refer to the source for more details.

At runtime, the device implements a small buffer for incoming event triplets (each one is stored as three 32-bit integers in a circular buffer), and raises its IRQ to signal them to the kernel.

When that happens, the kernel driver should use `IO_READ(READ)` to extract the 32-bit values from the device. Note that three `IO_READ()` calls are required to extract a single event triplet.

There are a few important notes here:

- The IRQ should not be raised `_before_` the kernel driver is started (otherwise the driver will be confused and ignore all events).

I.e. the emulator can buffer events before kernel initialization completes, but should only raise the IRQ, if needed, lazily. Currently this is done on the first `IO_READ(LEN)` following a `IO_WRITE(SET_PAGE, PAGE_ABSDATA)`.

- The IRQ is lowered by the device once all event values have been read, i.e. its buffer is empty.

However, on x86, if after an `IO_READ(READ)`, there are still values in the device's buffer, the IRQ should be lowered then re-raised immediately.

IX. Goldfish NAND device:

=====

Relevant files:

\$QEMU/hw/android/goldfish/nand.c
\$KERNEL/drivers/mtd/devices/goldfish_nand.c

Device properties:

```
817     Name: goldfish_nand
818     Id: -1
819     IrqCount: 1
820     I/O Registers:
821
822     This virtual device can provide access to one or more emulated NAND memory
823     banks [3] (each one being backed by a different host file in the current
824     implementation).
825
826     These are used to back the following virtual partition files:
827
828     - system.img
829     - data.img
830     - cache.img
831
832     TODO(digit): Complete this.
833
834
835     [3] http://en.wikipedia.org/wiki/Flash\_memory#NAND\_memories
836
837
838     X. Goldfish MMC device:
839     =====
840
841     Relevant files:
842     $QEMU/hw/android/goldfish/mmc.c
843     $KERNEL/drivers/mmc/host/goldfish.c
844
845     Device properties:
846     Name: goldfish_mmc
847     Id: -1
848     IrqCount: 1
849     I/O Registers:
850
851     Similar to the NAND device, but uses a different, higher-level interface
852     to access the emulated 'flash' memory. This is only used to access the
853     virtual SDCard device with the Android emulator.
854
855     TODO(digit): Complete this.
856
857
858     XIV. QEMU Pipe device:
859     =====
860
861     Relevant files:
862     $QEMU/hw/android/goldfish/pipe.c
863     $KERNEL/drivers/misc/qemupipe/qemu_pipe.c
864
865     Device properties:
866     Name: qemu_pipe
867     Id: -1
868     IrqCount: 1
869     I/O Registers:
870     0x00  COMMAND          W: Write to perform command (see below).
871     0x04  STATUS           R: Read status
```



```
872     0x08  CHANNEL      RW: Read or set current channel id.
873     0x0c  SIZE        RW: Read or set current buffer size.
874     0x10  ADDRESS     RW: Read or set current buffer physical address.
875     0x14  WAKES       R: Read wake flags.
876     0x18  PARAMS_ADDR_LOW RW: Read/set low bytes of parameters block address.
877     0x1c  PARAMS_ADDR_HIGH RW: Read/set high bytes of parameters block address.
878     0x20  ACCESS_PARAMS W: Perform access with parameter block.
```

```
879
```

```
880 This is a special device that is totally specific to QEMU, but allows guest
881 processes to communicate directly with the emulator with extremely high
882 performance. This is achieved by avoiding any in-kernel memory copies, relying
883 on the fact that QEMU can access guest memory at runtime (under proper
884 conditions controlled by the kernel).
```

```
885
```

```
886 Please refer to $QEMU/docs/ANDROID-QEMU-PIPE.TXT for full details on the
887 device's operations.
```

```
888
```

```
889
```

```
890 XIII. QEMU Trace device:
```

```
891 =====
```

```
892
```

```
893 Relevant files:
```

```
894     $QEMU/hw/android/goldfish/trace.c
895     $KERNEL/drivers/misc/qemutrace/qemu_trace.c
896     $KERNEL/drivers/misc/qemutrace/qemu_trace_sysfs.c
897     $KERNEL/fs/exec.c
898     $KERNEL/exit.c
899     $KERNEL/fork.c
900     $KERNEL/sched/core.c
901     $KERNEL/mm/mmap.c
```

```
902
```

```
903 Device properties:
```

```
904     Name: qemu_trace
905     Id: -1
906     IrqCount: 0
907     I/O Registers:
```

```
908
```

```
909 TODO(digit)
```

```
910
```